

上机作业A1 讲评

2022/9/28

刘通宇

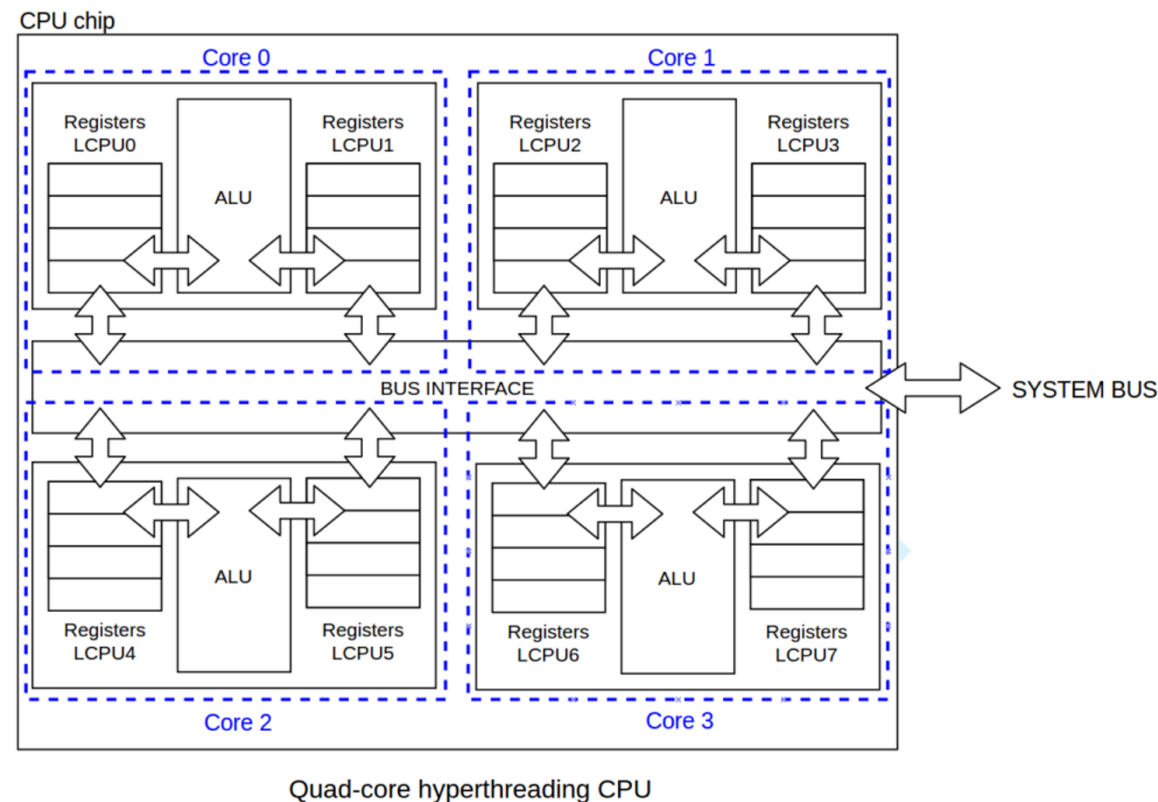
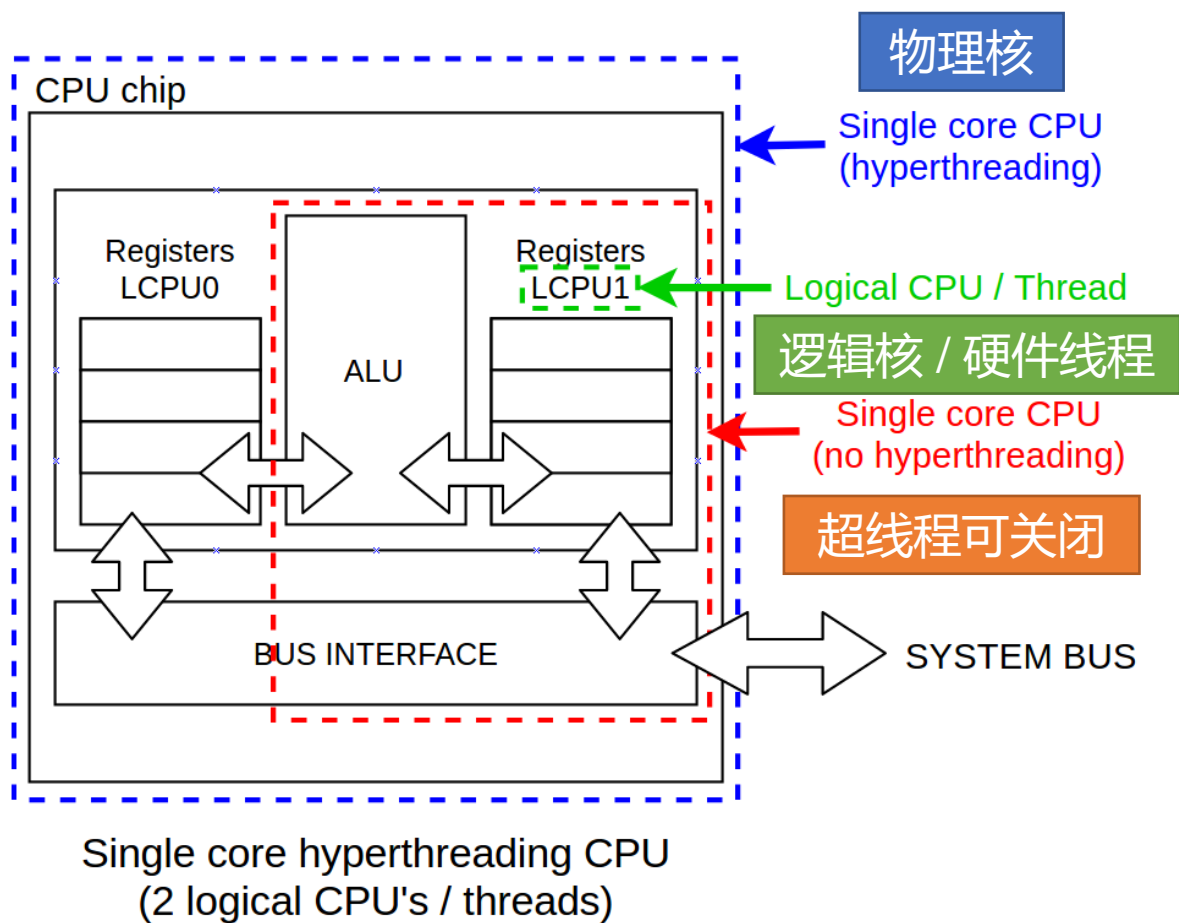
评分标准

- 上机作业A1的分数分为2部分：基础分数80%，额外分数20%
 - 基础分数：作业共2大题，16小题，每题5%，按回答是否正确给予客观评分
 - 额外分数：视作业完成态度、分析详尽程度与深入程度给予主观评分，最高20%
- 提示
 - 认真对待每一次Assignment和Project，每一次分数都是总评的一部分
 - 按照要求，准时地、完整地提交作业
 - 禁止抄袭

概念：指令集架构与微架构

- **指令集架构** (Instruction Set **Architecture**, ISA) 是计算机体系结构中与设计有关的部分，包含基本数据类型，指令集，寄存器，寻址模式，存储体系，中断，异常处理以及外部IO。
- **微架构** (Micro-**architecture**) 指的是一套用于执行指令集的处理器设计方法，使得指令集架构可以在处理器上被执行。
- 不同微架构的处理器可以共享一种指令集，例：Intel处理器与AMD处理器都是属于x86-64的指令集架构，但是两者在处理器的内部设计上存在本质区别。
- 不同代际 (generation) 的处理器，即使使用相同指令集，微架构层面上是有区别的，例：Intel第2代Xeon服务器芯片微架构代号是CascadeLake，第3代是IceLake。

Hyper-Threading 超线程




- Intel “四核八线程处理器”
 - 物理核：4个，逻辑核：8个
 - 操作系统能看到8个处理器供调度

类似Intel “超线程” 的技术

- 相似的技术

- Intel® Hyper-Threading, 超线程, 简称HT
- AMD® Simultaneous Multithreading, 同步多线程, 简称SMT

- 思考: HT或SMT带来的性能提升?



lscpu的输出

Server with Intel processors

```
(base) tongyu@solesystem:~$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:             Little Endian
Address sizes:          46 bits physical, 48 bits virtual
CPU(s):                 80
On-line CPU(s) list:   0-79
Thread(s) per core:     2
Core(s) per socket:     20
Socket(s):              2
NUMA node(s):          2
Vendor ID:              GenuineIntel
CPU family:             6
Model:                 85
Model name:             Intel(R) Xeon(R) Gold 5218R CPU @ 2.10GHz
Stepping:               7
.....
```

这里的 “Architecture” 指的是指令集架构

“Thread(s) per core” 每一个物理核对应的硬件线程个数

Server with Kunpeng processors

```
(base) tongyu@taishan-200:~$ lscpu
Architecture:          aarch64
CPU op-mode(s):        64-bit
Byte Order:             Little Endian
CPU(s):                96
On-line CPU(s) list:   0-95
Vendor ID:              HiSilicon
Model name:             Kunpeng-920
Model:                 0
Thread(s) per core:     1
Core(s) per socket:     48
Socket(s):              2
Stepping:               0x1
.....
```

信息来源于处理器厂商提供的信息，这里的 “Model” 是厂商内部的型号标识，具体含义需要参考相应的架构文档。

CPUID命令

- 对于x86-64指令集架构，CPUID是其指令集下的一条机器指令，这条指令通过EAX的取值，判断需要取得的信息，再通过EAX、EBX、ECX、EDX作为返回值，得到相应信息，这些取值和返回信息需要参考硬件生产商的相关文档。

- 对Intel的处理器按照DisplayFamily_DisplayModel的模式查找，06_9EH所在的条目

[CPUID - Wikipedia](#)

`man cpuid`: [cpuid\(1\): Dump CPUID info for each CPU - Linux man page \(die.net\)](#)

Intel® 64 and IA-32 Architectures Software Developer's Manuals Volume 4 Chapter 2 Model Specific Registers (MSRs)

```
CPU 0:
vendor_id = "GenuineIntel"
version information (1/eax):
  processor type = primary processor (0)
  family = 0x6 (6)
  model = 0xe (14)
  stepping id = 0x9 (9)
  extended family = 0x0 (0)
  extended model = 0x9 (9)
  (family synth) = 0x6 (6)
  (model synth) = 0x9e (158)
  (simple synth) = Intel Core (unknown type) (Kaby Lake /
Coffee Lake) [Kaby Lake] {Skylake}, 14nm
.....
```

Table 2-1. CPUID Signature Values of DisplayFamily_DisplayModel

DisplayFamily_DisplayModel	Processor Families/Processor Number Series
06_85H	Intel® Xeon Phi™ Processor 7215, 7285, 7295 Series based on Knights Mill microarchitecture
06_57H	Intel® Xeon Phi™ Processor 3200, 5200, 7200 Series based on Knights Landing microarchitecture
06_8CH, 06_8DH	11th generation Intel® Core™ processors based on Tiger Lake microarchitecture
06_7DH, 06_7EH	10th generation Intel® Core™ processors based on Ice Lake microarchitecture
06_A5H, 06_A6H	10th generation Intel® Core™ processors based on Comet Lake microarchitecture
06_66H	Intel® Core™ processors based on Cannon Lake microarchitecture
06_8EH, 06_9EH	7th generation Intel® Core™ processors based on Kaby Lake microarchitecture, 8th and 9th generation Intel® Core™ processors based on Coffee Lake microarchitecture, Intel® Xeon® E processors based on Coffee Lake microarchitecture
06_6AH, 06_6CH	3rd generation Intel® Xeon® Processor Scalable Family based on Ice Lake microarchitecture
06_55H	Intel® Xeon® Processor Scalable Family based on Skylake microarchitecture, 2nd generation Intel® Xeon® Processor Scalable Family based on Cascade Lake product, and 3rd generation Intel® Xeon® Processor Scalable Family based on Cooper Lake product
06_4EH, 06_5EH	6th generation Intel Core processors and Intel Xeon processor E3-1500m v5 product family and E3-1200 v5 product family based on Skylake microarchitecture
06_5FH	Intel Xeon processor D-1500 product family based on Broadwell microarchitecture

Linux操作系统的内核版本与发行版本

- 内核 (Kernel) : Linux的底层和核心部分, 是硬件与软件之间的中间层, 充当底层的驱动程序, 负责将可用的资源分配到各个进程。
- 发行版本 (Distribution) : 在内核的基础上, 开发不同应用程序, 组成的一个完整的操作系统, 用户可以直接使用, 例如RedHat, Debian, Ubuntu等。

uname -a的输出

```
(base) tongyu@taishan-200:~$ uname -a
Linux taishan-200 5.15.0-41-generic #44-Ubuntu
SMP Thu Jun 23 11:20:13 UTC 2022 aarch64
aarch64 aarch64 GNU/Linux
```

- 操作系统名称 (-s选项)
- 计算机名称 (-n选项)
- **操作系统内核版本 (-r选项)**
- 发行版版本与时间 (-v选项)
- 指令集架构 (-m选项)

* SMP: Symmetric Multi-Processor 对称多处理器

numactl -H与dmidecode的输出

Note: 在最新的WSL2上, 这两条命令无法导出机器的信息

```
(base) tongyu@taishan-200:~$ numactl -H
available: 4 nodes (0-3)
node 0 cpus: 0 1 2 3 4 5 6 7 8 9 10 11 12 13
14 15 16 17 18 19 20 21 22 23
node 0 size: 47883 MB
node 0 free: 257 MB
.....
node distances:
node  0  1  2  3
 0:  10  16  32  33
 1:  16  10  25  32
 2:  32  25  10  16
 3:  33  32  16  10
.....
```

```
(base) root@taishan-200:~# dmidecode | grep -A16 "Memory Device"
Memory Device
    Array Handle: 0x0006
    Error Information Handle: Not Provided
    Total Width: 72 bits
    Data Width: 64 bits
    Size: 16 GB
    Form Factor: DIMM
    Set: None
    Locator: DIMM000 J27
    Bank Locator: SOCKET 0 CHANNEL 0 DIMM 0
    Type: DDR4
    Type Detail: Synchronous Registered (Buffered)
    Speed: 2933 MT/s
    Manufacturer: Hynix
    Serial Number: 20E7B5EE
    Asset Tag: 2011
    Part Number: HMA82GR7CJR4N-WM
```

Question: 这里的GB, 指的是 2^{20} Byte还是 10^6 Byte?

Python2和Python3

- Python3在设计的时候没有考虑向下兼容，因此使用Python2的程序无法直接使用Python3的解释器执行。

Python2: `print "Hello, world"`

Python3: `print("Hello, world")`

- 我已经安装了Python3了，怎么办？

1. Download source code

```
$ wget https://www.python.org/ftp/python/2.7.9/Python-2.7.9.tgz
```

2. Unzip, compile to install

```
$ tar -zxvf Python-2.7.9.tgz
```

```
$ cd Python-2.7.9
```

```
$ ./configure --prefix=~/.local/python-2.7.9
```

```
$ make
```

```
$ make install
```

3. Create a symbolic link

```
$ ln -s ~/.local/python-2.7.9/bin/python ~/.local/python27
```

4. Add to \$PATH

```
export PATH=~/.local/bin:$PATH
```

5. Run

```
$ python27 helloworld.py
```